

# T-DeLearn PW10

---

- Romain Capocasale
- Jean Demeusy

## Exercice 1

A

Summary of sequential architecture :

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
activation (Activation)	(None, 32, 32, 32)	0
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
activation_1 (Activation)	(None, 32, 32, 32)	0
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 32)	9248
activation_2 (Activation)	(None, 16, 16, 32)	0
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 32)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 10)	20490
Total params: 39,882		
Trainable params: 39,882		
Non-trainable params: 0		

Test loss: 0.8561450242996216

Test accuracy: 0.714900016784668

Score of sequential :

Summary of functional architecture :

```
visible = Input(shape=X_train.shape[1:])
conv1 = Conv2D(32, kernel_size=(3,3), padding="same", activation="relu")(visible)
conv2 = Conv2D(32, kernel_size=(3,3), padding="same", activation="relu")(conv1)
pool1 = MaxPooling2D(2)(conv2)
conv3 = Conv2D(32, kernel_size=(3,3), padding="same", activation="relu")(pool1)
pool2 = MaxPooling2D(2)(conv3)
flat = Flatten()(pool2)
output = Dense(10, activation="softmax")(flat)

model = Model(inputs=visible, outputs=output)
model.summary()
```

Model: "model\_2"

Layer (type)	Output Shape	Param #
input_7 (InputLayer)	(None, 32, 32, 3)	0
conv2d_13 (Conv2D)	(None, 32, 32, 32)	896
conv2d_14 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d_9 (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_15 (Conv2D)	(None, 16, 16, 32)	9248
max_pooling2d_10 (MaxPooling2D)	(None, 8, 8, 32)	0
flatten_5 (Flatten)	(None, 2048)	0
dense_5 (Dense)	(None, 10)	20490
Total params: 39,882		
Trainable params: 39,882		
Non-trainable params: 0		

Test loss: 0.9130173603057862

Test accuracy: 0.7125999927520752

Score of functional :

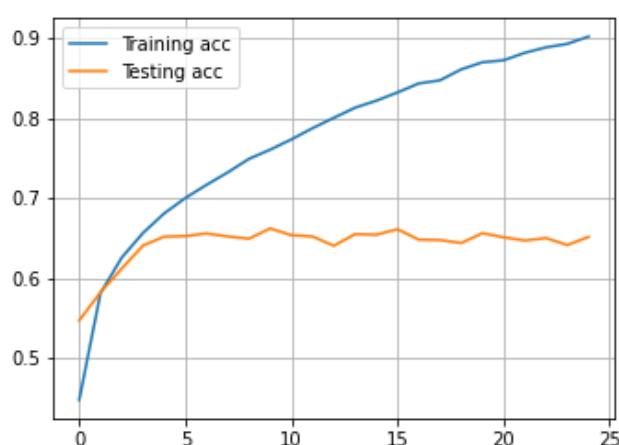
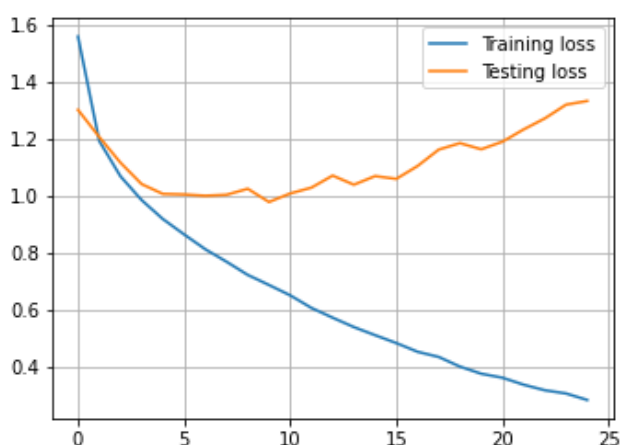
We can see that, the two models have the same number of parameters and the scores are slightly almost similar.

B

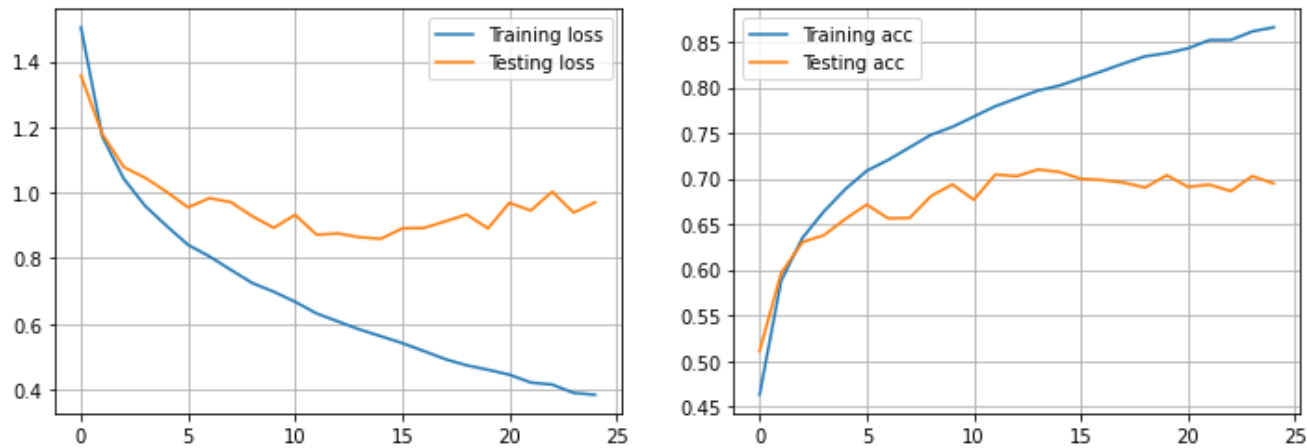
Model	Architecture description	Acc. train	Acc. test
1 - multiple path	BRANCH1: CONV(D=32, w=h=3, S=1, P=same, act=relu)- DROPOUT(0.2)-MAXPOOL(S=1, size=2)-FLATTEN; BRANCH2: CONV(D=32, w=h=6, S=1, P=same, act=relu)- DROPOUT(0.2)-MAXPOOL(S=1, size=2)-FLATTEN; CONCAT(BRANCH1, BRANCH2); DENSE(100, act=relu); DENSE(10, softmax)	0.9020	0.6517

Model	Architecture description	Acc. train	Acc. test
2 - multiple features	BRANCH1:CONV(D=32, w=h=3, S=1, P=same, act=relu)- DROPOUT(0.2)-MAXPOOL(S=1, size=2)-FLATTEN; BRANCH2(BRANCH1):CONV(D=32, w=h=3, S=1, P=same, act=relu)- DROPOUT(0.2)-MAXPOOL(S=1, size=2)-FLATTEN; BRANCH3(BRANCH2):CONV(D=32, w=h=3, S=1, P=same, act=relu)- DROPOUT(0.2)-MAXPOOL(S=1, size=2)-FLATTEN; CONCAT(BRANCH1, BRANCH2, BRANCH3); DENSE(100, act=relu); DENSE(10, softmax)	0.8658	0.6945
3 - multiple features	BRANCH1:CONV(D=32, w=h=3, S=1, P=same, act=relu)- DROPOUT(0.2)-MAXPOOL(S=1, size=2)-FLATTEN; BRANCH2(BRANCH1):CONV(D=32, w=h=3, S=1, P=same, act=relu)- DROPOUT(0.2)-MAXPOOL(S=1, size=2)-FLATTEN; BRANCH3(BRANCH2):CONV(D=32, w=h=3, S=1, P=same, act=relu)- DROPOUT(0.2)-MAXPOOL(S=1, size=2)-FLATTEN; CONCAT(BRANCH1, BRANCH2, BRANCH3); DENSE(64, act=relu); DROPOUT(0.2); DENSE(10, softmax)	0.8216	0.7153
4 - multiple path	BRANCH1:CONV(D=32, w=h=3, S=1, P=same, act=relu)- DROPOUT(0.2)-MAXPOOL(S=1, size=2)-CONV(D=64, w=h=3, S=1, P=same, act=relu)-DROPOUT(0.2)-MAXPOOL(S=1, size=2)-FLATTEN; BRANCH2(BRANCH1):CONV(D=32, w=h=3, S=1, P=same, act=relu)- DROPOUT(0.2)-MAXPOOL(S=1, size=2)-CONV(D=64, w=h=6, S=1, P=same, act=relu)-DROPOUT(0.2)-MAXPOOL(S=1, size=2)-FLATTEN; CONCAT(BRANCH1, BRANCH2); DENSE(50, act=relu); DROPOUT(0.2); DENSE(10, softmax)	0.6765	0.6891

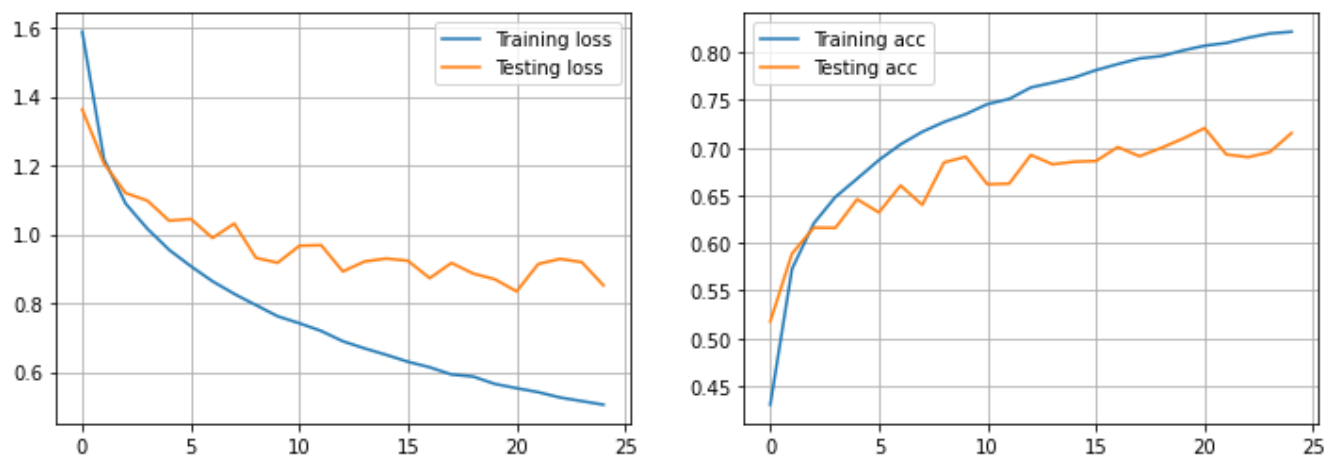
Model 1 Loss/Accs :



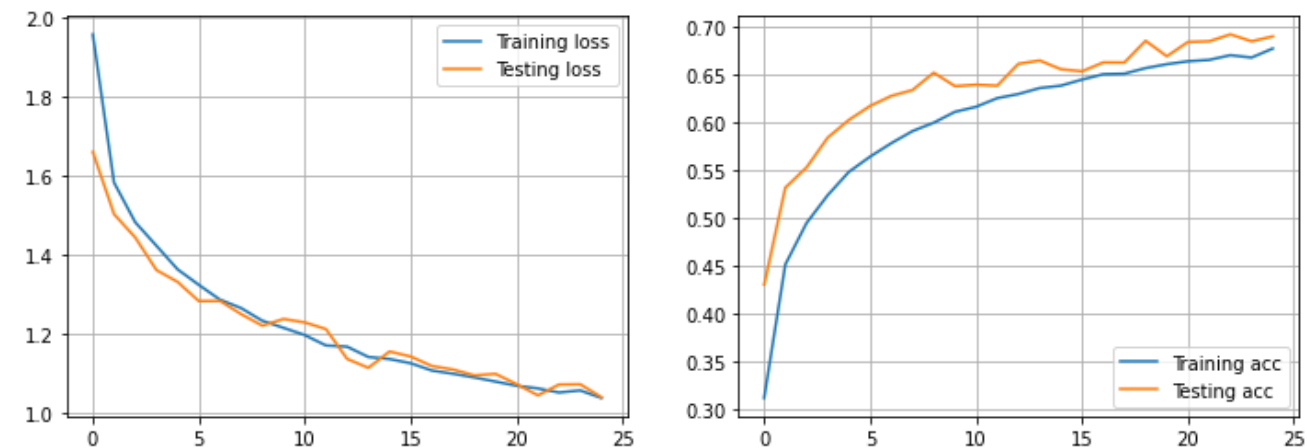
Model 2 Loss/Accs :



Model 3 Loss/Accs :



Model 4 Loss/Accs :



Discussion

It can be seen that the model with the best accuracy is model number 3 with multiple features. We also see that the first 3 models overfit, this problem can be avoided by using data augmentation which has not been implemented. The 4th model overfits the least and seems the most stable.

We notice with this lab and the 2 previous ones that the accuracy of the models is around 0.70% and cannot be significantly exceeded despite all attempts.

# Exercice 2

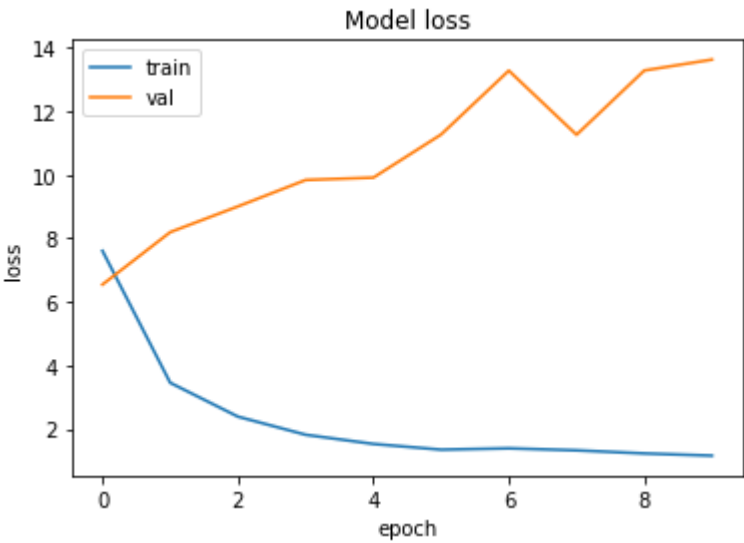
The base model chosen is MobileNet.

## Architecture 1

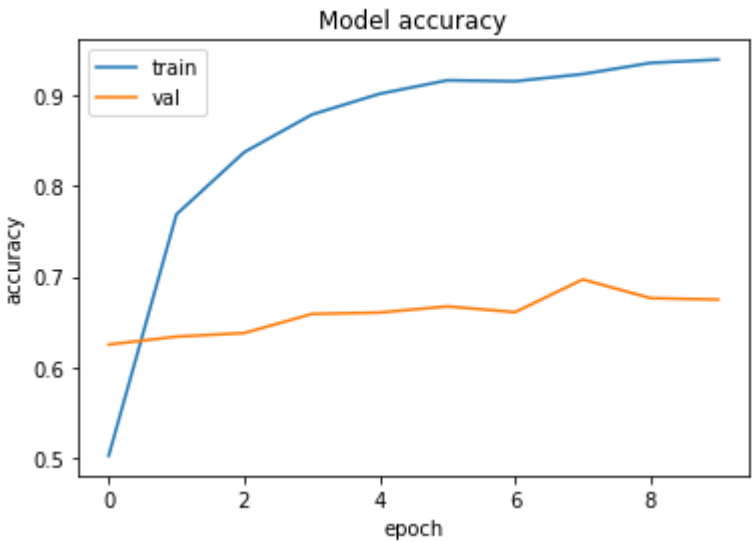
Hyperparameters:

- Learning Rate : 0.001
- Batch Size = 128
- Epoch = 20

Model	Architecture description	Acc. train	Acc. test
1	MobileNet - Flatten - Dropout(0.5) - Dense(20)	0.94	0.68



Loss architecture 1:



Accuracy architecture 1:

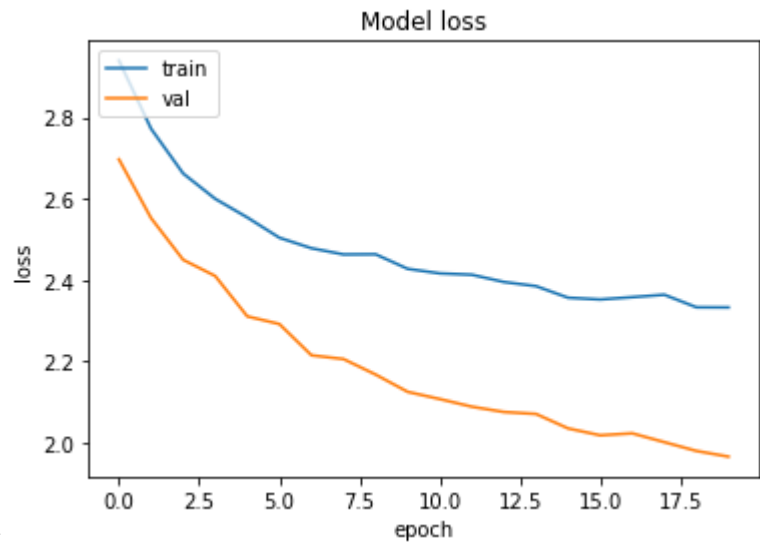
We can see that despite the performance of the model, it overfits a lot and will not be retained for the fine-tuning phase.

## Architecture 2

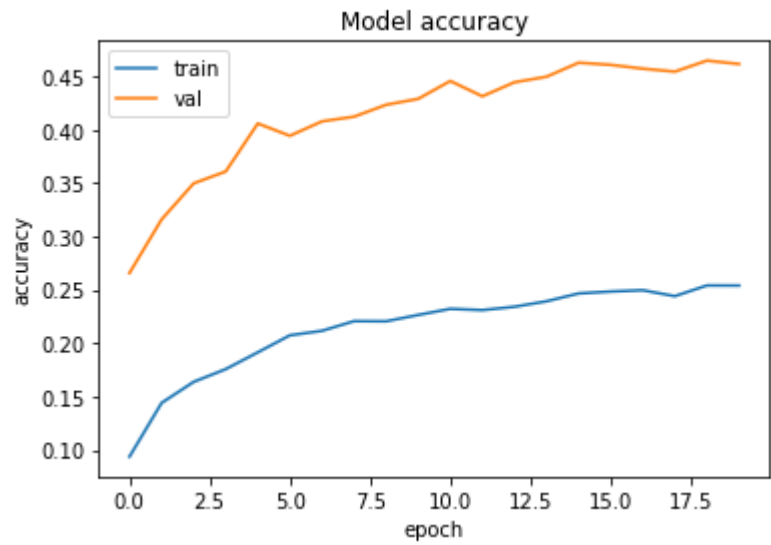
Hyperparameters:

- Learning Rate : 0.001
- Batch Size = 128
- Epoch = 20

Model	Architecture description	Acc. train	Acc. test
2	MobileNet - Flatten - Dropout(0.5) - Dense(32) - Dropout(0.5) - Dense(20)		



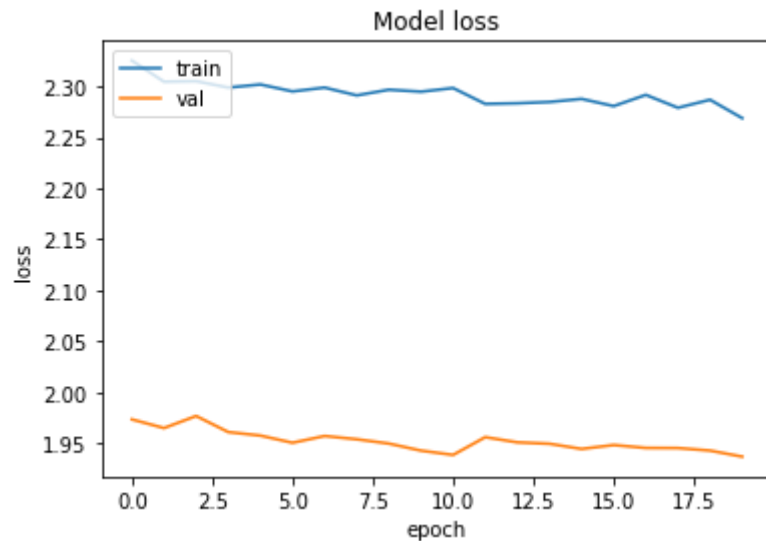
Loss architecture 2:



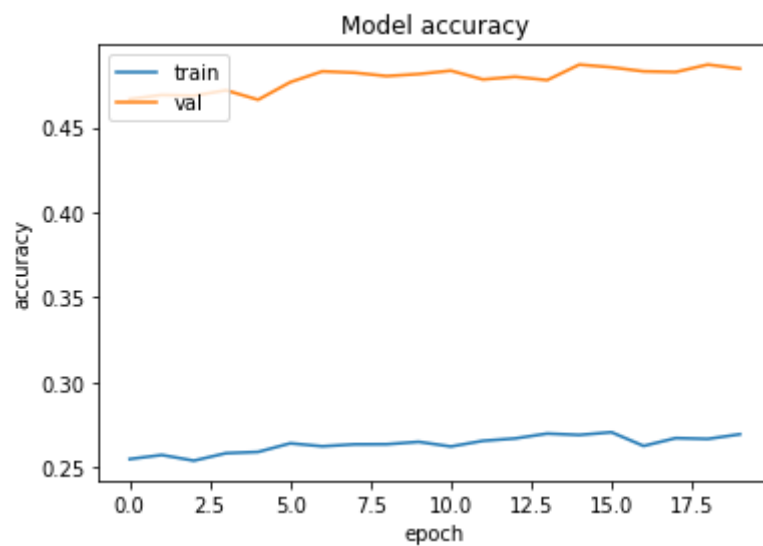
Accuracy architecture 2:

Hyperparameters fine tuning:

- Learning Rate : 0.00000001
- Batch Size = 128
- Epoch = 20



Loss fine-tuning architecture 2:



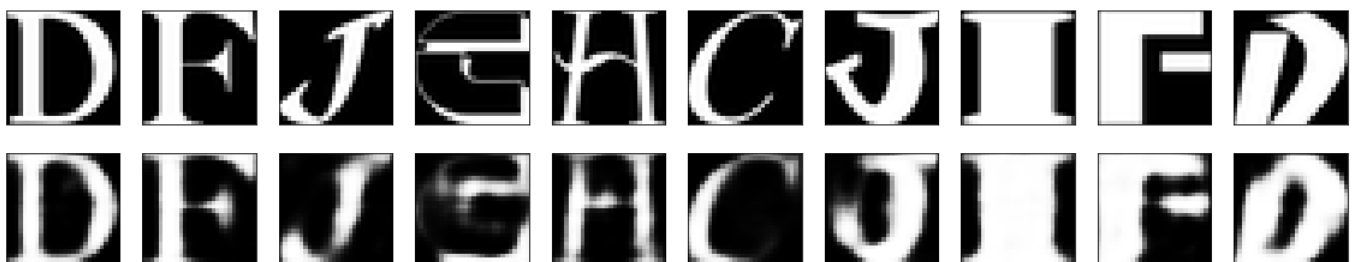
Accuracy fine-tuning architecture 2:

## Discussion

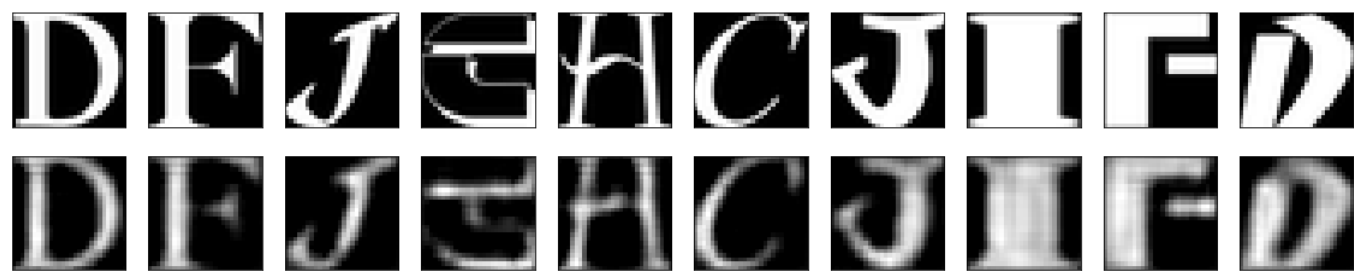
We can see that despite the fact that the first architecture gives better results it overfits a lot. The second architecture gives less good results but overfits less. This second model is therefore more stable and has been retained to perform the finetuning. We can see that the finetuning has only slightly improved the score of the model. The accuracy has gone from about **0.4614** to **0.4845**.

## Exercise 3

Input and output of the autoencoder:



Prediction with autoencodeur:



f) We could train a classifier on the input data. The model obtained is then saved, and used with the images obtained with the decoder. Then the performance could be estimated with the evaluation of recall, f1scor or precision.

g) With the SVM, we get a calssification accuracy of 90%. SVM results:

Classification report for classifier SVC(cache_size=7000, kernel='linear'):				
	precision	recall	f1-score	support
0	0.91	0.90	0.90	1000
1	0.90	0.86	0.88	1000
2	0.91	0.92	0.91	1000
3	0.91	0.90	0.91	1000
4	0.90	0.88	0.89	1000
5	0.90	0.94	0.92	1000
6	0.88	0.88	0.88	1000
7	0.91	0.89	0.90	1000
8	0.86	0.86	0.86	1000
9	0.87	0.93	0.90	1000
accuracy			0.90	10000
macro avg	0.90	0.90	0.90	10000
weighted avg	0.90	0.90	0.90	10000

Denoising with autoencodeur (noise\_level = 0.3):

